


Your position: RECKLESS

producing value, accumulating risk



What this position means

Real metrics, genuine productivity, measurable value — and near-zero enforced governance. This is where most high-impact shadow AI lives, and it feels like success from the inside because the value is real. The debt is also real and simply not yet visible. The primary risk is not that a system fails today; it is that when one does — and drift makes failure a matter of when — there is no audit trail, no rollback, no named owner, no incident process. The technical failure is recoverable. The organisational overcorrection that follows is usually more expensive than governance would have been: legal freezes deployments, security demands central review of everything, and the AI programme loses a year.

Direction of travel: Reckless → Governed

Primarily an inventory-and-ownership problem, not a technology problem. The estate is producing value through systems the organisation cannot see, bound, or explain. The work is to make the value governable without killing it.

The first three moves

- **1. Inventory the shadow estate.** Run multi-signal discovery — SSO/OAuth logs, expense data, SaaS and cloud usage, repository scans, interviews. A register fed only by project approvals is the reason you are here. Include vendor AI features activated inside platforms you already own.
- **2. Name owners where the benefit sits.** For the highest-value systems first: a named individual in the business unit that profits from the system, not a committee, not the vendor, not the team that happened to build it.
- **3. Apply the four minimums to the top of the value list.** Named owner, technically enforced data boundary, monitoring baseline with a review cadence, tested fallback. Entry conditions, not finishing touches — and highest-Value-Velocity systems move first, because they combine the most value with the most exposure.

The characteristic obstacle

The conversation with the team that built the shadow system. They moved fast because moving fast was rewarded, and bringing them into governance reads as punishment. Two durable answers: change what is rewarded, or make the governed path faster than the ungoverned one — reusable access routes, pre-approved patterns, standard logging — so the incentive problem solves itself.

Watch for: Do not overcorrect. The value in your Reckless cluster is real; freezing everything converts governance debt into opportunity cost and teaches the organisation that governance kills value. Sequence by value-at-risk instead.

What to measure

- Share of high-value systems with a named owner and an enforced data boundary (target: the top of the value list first, not uniform coverage).
- Time-to-governed per system — falling with each migration if governance components are being reused.
- Count of unmanaged data connections closed.

Rescoring — the instrument itself

Your questionnaire result is a screening across the estate on average. The instrument proper is applied per system.

Governance Discipline (0–4):

0	No owner, no documentation, no monitoring. The system exists; nobody is accountable.
1	Owner named. No enforced controls, no monitoring baseline. Governance is a name on a list.
2	Policy exists and is <i>partially</i> enforced. Monitoring manual or intermittent. (A policy nobody can show is enforced scores 1, not 2.)
3	Policy enforced consistently. Active monitoring with a review cadence. Incident escalation documented <i>and tested</i> .
4	Every governance claim is demonstrable from the system's own telemetry. The system could face a regulator with evidence behind every claim.

Scores 0–1 plot Low; 3–4 High; a 2 is the transition zone — err low. **Value Velocity** is one question: can we attribute and quantify this system's KPI contribution, with a defensible number and methodology — yes or no?

The exercise, in five steps

Capture leadership's belief first and do not let anyone revise it. Build the real inventory from multiple signals — not from project registrations. Score every system with the rubric above, erring low where evidence is missing (missing evidence is the finding). Plot the map. Measure the distance between the belief and the data — that gap is the headline of the board report. Repeat quarterly: the Quadrant is a snapshot, and the direction of travel matters more than any single placement.

Status of this instrument

A working diagnostic from the field research behind *Govern or Fail: Why Enterprise AI Fails After the Demo* — deliberately published here rather than in the book, which stops where validation stops. The book is the diagnosis; this is the field kit. Use it, and where it breaks, tell me: governorfail.com. *Govern or Fail* launches September 2026.