

Ten pressure tests for enterprise AI governance

*A working index of the diagnostic instruments from **Govern or Fail: A Diagnostic Field Report on Enterprise AI Governance***

Most enterprise AI failure is not caused by weak models. It is caused by weak operating conditions — the gap between the governance an organisation can describe and the governance it can demonstrate. These ten instruments make that gap visible before an incident does. They are not a methodology and not a sequence: each one pressure-tests a claim someone in the organisation is already making — that the words are aligned, that the pilot proved something, that the vendor system is auditable, that certification means control. Use whichever instrument matches the claim currently on the table.

■ 1. The Language Questions (Chapter 3)

Five questions that test whether the people funding, building, buying, and approving an AI system mean the same thing by the words that define it: what “agent” means in this deployment, what “production-ready” covers, who owns the system after go-live, what would signal behavioural drift, and what the regulatory classification is.

Use before serious budget is committed.

■ 2. The Portfolio Lens (Chapter 5)

Two questions applied to every system in the estate: how much governance discipline does it actually have, and how much production value does it actually generate? The tension between the answers is where the real portfolio becomes visible.

Use when the official inventory feels more complete than the estate.

■ 3. The Operating Test (Chapter 6)

When the next AI system moves toward production, ask what the team must build from scratch. If the answer is everything — access route, logging, monitoring, evidence, review design, incident procedure — governance is not yet reusable. It is another project.

Use to test whether governance is compounding or repeating.

■ 4. The Production Threshold (Chapter 7)

The questions that separate an experiment from a production candidate: is the data path mapped, the owner named, the use case risk-classified, the monitoring and evidence requirements understood, the cost structure estimated, the fallback defined — and can human oversight operate at the volume the system will create?

Use before a pilot is treated as proof of scalability.

■ 5. The Vendor Questions (Chapter 8)

Three questions that force implementation details into view: what is the smallest unit of behaviour we can reconstruct, where is the boundary enforced, and how do we recover from a systematic failure?

Use before buying or approving a system that will act inside a material process.

■ 6. The Executive Translation (Chapter 9)

Five questions a technical briefing should answer: what source was used, what could the system access or do, what happened between components, where was policy enforced, and can the decision path be reconstructed?

Use whenever an architecture discussion starts producing acronyms instead of answers.

■ 7. The Customer Challenge Test (Chapter 10)

Choose one recent AI-supported decision. Assume it is disputed. Work through the questions an investigator would ask. If the answer requires a special investigation across logs, exports, emails, and memory, the organisation has governance documents without operational governance.

The harshest instrument on this list — and the most useful.

■ 8. The Portfolio Metrics (Chapter 12)

Six illustrative measures that connect governance, value, and velocity: estate visibility, named ownership coverage, runtime evidence coverage, reuse rate, time to incident response, and value attribution.

Use to move the board conversation beyond adoption counts.

■ 9. The Three Pressure Tests (Chapter 12)

Where does the organisation actually sit on the gap between governance described and governance demonstrated? What does the current trajectory look like compounded to 2030? What architectural decision is being made — or deferred — this quarter?

Use to close an executive session.

■ 10. The Four Prerequisites (Afterword)

Discovery, ownership, boundaries, route. Not an operating model — the gating conditions any operating model would have to build on.

Use to decide what must exist before any framework is worth adopting.

Why there is no framework in this document

None of these instruments tells the organisation what to build. Together they answer a narrower question: whether the governance the organisation describes is governance it can demonstrate. The operational specifics — the architecture, the operating rhythm, the build-versus-buy decisions — deserve more validation than a checklist can provide. That is deliberate, and it is the subject of the work that follows.

The book

Govern or Fail: A Diagnostic Field Report on Enterprise AI Governance — a diagnostic field report on enterprise AI governance. Twelve chapters on why AI initiatives collapse between the demo and production, built on the strongest available evidence and honest about where the evidence stops. Available September 2026. Updates, source notes, and new essays at governorfail.com.